

## 1. Executive Summary

Enterprises are adopting artificial intelligence faster than they can govern it. Every time a company sends data to an AI model, it risks leaking sensitive customer or patient information, breaching regulations such as GDPR or HIPAA, and failing the audits that large customers now demand before they will buy. Today, most companies manage this risk with manual reviews, spreadsheets, and hope.

**AuthClaw closes that gap.** It is a software product that sits between a company's applications and the AI models they use, acting as an automated safety checkpoint. It removes sensitive data in real time, enforces the company's rules, fixes compliance problems with human sign-off, and keeps a tamper-proof record that can be handed straight to an auditor.

This briefing explains — in non-technical terms — what AuthClaw is, why the timing is right, how large the opportunity is, who buys it, what it will cost to build, and the decision we are asking the Board to approve.

### The ask, in one line

Approve a four-stage build of the AuthClaw product — roughly **2,560 engineering hours** (an indicative engineering investment of about **US\$123,000**) — to reach a sellable, audit-ready first version.

Third-party costs (security testing, SOC 2 audit) and go-to-market are additional and are flagged separately in Section 10.

## 2. The Problem We Are Solving

Three forces are colliding inside almost every mid-to-large company right now:

### 2.1 AI is being adopted everywhere, quickly

Teams across customer support, finance, legal, and operations are wiring AI models into their daily work. Most of this happens without a central safety layer, so the company often cannot say exactly what data is being sent to which AI provider.

### 2.2 Regulators are raising the stakes

Data-protection laws (GDPR, HIPAA) and security standards (SOC 2) already carry real penalties. The EU's new AI Act goes further: non-compliance can cost up to **€35 million or 7% of global annual turnover**, whichever is higher. For a large enterprise, that is a board-level financial risk, not an IT footnote.

### 2.3 Customers now audit their suppliers

Large buyers increasingly refuse to sign contracts until a supplier can prove its security and compliance posture. Today that proof is assembled by hand for every deal, slowing sales and tying up senior staff.

### The cost of doing nothing

Sensitive data can leak to external AI providers without anyone noticing until it is too late.

A single regulatory finding or breach can trigger fines, lost customers, and reputational damage. Sales cycles stall while teams manually produce compliance evidence deal-by-deal.

### 3. What AuthClaw Is (in plain terms)

AuthClaw is delivered as three connected capabilities. A useful analogy: it works like a combination of a security scanner at an airport, a knowledgeable compliance officer, and a black-box flight recorder — applied to a company's use of AI.

| Capability                            | What it does, in plain English                                                                                                                                                                                                                              | Business benefit                                                     |
|---------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------|
| 1. The Safety Checkpoint ("Gateway")  | Every message going to or from an AI model passes through AuthClaw first. It automatically detects and removes sensitive information (names, health data, card numbers) and blocks anything that breaks company rules — before it ever leaves the building. | Prevents data leaks and policy breaches in real time.                |
| 2. The Compliance Assistant ("Agent") | An AI assistant that scans the company's systems, explains in clear language where it falls short of GDPR / HIPAA / SOC 2, and proposes fixes. Crucially, it never makes a risky change on its own — a human must approve each one, with a security check.  | Turns weeks of manual compliance work into guided, approved actions. |
| 3. The Audit Recorder ("Audit Trail") | Every action, decision, and approval is written into a tamper-proof record that cannot be secretly altered. This can be exported as verifiable proof for auditors and customers.                                                                            | Replaces manual evidence-gathering with one trustworthy report.      |

#### The safeguard that matters most to a Board

AuthClaw can recommend and prepare fixes, but it cannot execute anything risky without an explicit human approval backed by a security check, and that approval automatically expires if not acted on quickly.

In short: the machine does the heavy lifting; a person always stays in control of consequential decisions.

### 4. Why Now

AuthClaw is not an idea ahead of its market. The timing is being driven by forces outside our control, all pointing the same way:

- **Regulation has teeth.** The EU AI Act began phasing in through 2025, with severe penalties. Other regions are following. Compliance is becoming a mandatory budget line, not an optional one.
- **AI adoption is mainstream.** The question for enterprises is no longer whether to use AI, but how to use it safely — exactly the gap AuthClaw fills.
- **Buyers are paying for trust.** Insurers are now linking premiums to certified governance, and customers gate deals on proof of compliance. Trust has become a purchasing criterion.
- **The category is funded and proven.** Established companies already sell adjacent pieces of this puzzle (see Section 6), which de-risks the market: we are entering a validated category, not inventing one.

## 5. Market Opportunity

AI governance is an early but fast-growing market. Independent analysts disagree on the exact size — a normal sign of a young category — but they agree on the shape: small today, expanding rapidly, multi-billion-dollar within a few years.

| Analyst (2026)           | 2026 size (approx.)      | Growth (CAGR) | Outlook                  |
|--------------------------|--------------------------|---------------|--------------------------|
| Research & Markets       | ~US\$0.6 billion         | ~44%          | ~US\$2.6 billion by 2030 |
| Mordor Intelligence      | ~US\$0.44 billion        | ~28%          | ~US\$1.5 billion by 2031 |
| Coherent Market Insights | ~US\$0.4 billion (2025)  | ~47%          | ~US\$6.0 billion by 2032 |
| Grand View Research      | ~US\$0.31 billion (2025) | ~36%          | ~US\$3.6 billion by 2033 |

Figures are independent third-party estimates as of early 2026 and are indicative, not guarantees. The consistent message: double-digit-to-~45% annual growth and a multi-billion-dollar market within roughly five years.

The fastest-growing slices are exactly where AuthClaw is strong: healthcare and life sciences, financial services, and any organisation deploying generative AI that must prove it is safe. Spending starts with large enterprises but is broadening to mid-sized companies as cloud-based tools lower the barrier.

## 6. Competitive Landscape & Our Wedge

Three well-known companies illustrate the category — and, importantly, each solves only one part of the problem. They are the reference points we benchmarked the product against.

| Company  | What they do well                                                                                                                                    | What they do NOT do                                                    |
|----------|------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------|
| Vanta    | Automates compliance certifications (SOC 2, ISO, GDPR) and produces a customer-facing trust page. A category leader with a very large customer base. | Does not sit in the live AI traffic to stop data leaks as they happen. |
| Securiti | Discovers and governs sensitive data across an organisation.                                                                                         | Not focused on intercepting and fixing AI interactions in real time.   |
| Giskard  | Tests AI models for weaknesses such as harmful answers and prompt manipulation.                                                                      | A testing tool, not a live gateway or a compliance-evidence system.    |

### AuthClaw's wedge: the runtime layer

Vanta proves you were compliant. Giskard tests whether your AI is safe. Securiti maps your data. AuthClaw is the live layer that **enforces safety in the moment, fixes gaps with human approval, and records the proof** — combining the strongest parts of all three into a single product positioned at the point where AI is actually used.

## 7. Who Buys AuthClaw, and Why

| Buyer segment                                          | Why they need it                                                          | Who signs off                  |
|--------------------------------------------------------|---------------------------------------------------------------------------|--------------------------------|
| Regulated enterprises (healthcare, finance, insurance) | Heaviest penalties and audit scrutiny; cannot afford an AI data leak.     | CISO / Chief Risk / Compliance |
| AI-first software companies                            | Must prove to their own customers that their AI is safe to unblock sales. | CTO / Head of Security         |

| Buyer segment               | Why they need it                                                       | Who signs off |
|-----------------------------|------------------------------------------------------------------------|---------------|
| Mid-market firms scaling AI | Adopting AI fast without a dedicated governance team; need automation. | CTO / COO     |

The common thread: the buyer is someone personally accountable for risk and audits. For them, AuthClaw is not a cost — it is insurance that also accelerates revenue.

## 8. Business Value & Outcomes

AuthClaw creates value on both sides of the ledger — it reduces risk and it helps customers earn revenue faster.

| Outcome                 | How AuthClaw delivers it                                                           |
|-------------------------|------------------------------------------------------------------------------------|
| Avoid fines & breaches  | Stops sensitive data leaving and enforces rules before any harm occurs.            |
| Win deals faster        | Provides instant, verifiable proof of compliance, removing a common sales blocker. |
| Cut compliance labour   | Replaces manual evidence-gathering with automated, exportable reports.             |
| Stay continuously ready | Monitors compliance posture in real time, not just once a year at audit.           |
| Keep humans in control  | Every consequential change requires explicit, secure human approval.               |

## 9. Build Roadmap — in Business Terms

The product is built in four stages. Each stage delivers something the business can show to customers and investors, so value compounds rather than arriving only at the end.

| Stage                        | What gets built (plain English)                                                                                                    | What we can show after it                                                           | Effort  |
|------------------------------|------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------|---------|
| 1. Foundation                | The secure plumbing: the safety checkpoint passing AI traffic, strict separation between customers, and the tamper-proof recorder. | A working demo of AI traffic flowing safely through AuthClaw.                       | 640 hrs |
| 2. Intelligence & Guardrails | Real-time removal of sensitive data, rule enforcement, and the AI compliance assistant with human-approval safeguards.             | A live demo of AuthClaw catching a leak and an assistant proposing an approved fix. | 640 hrs |
| 3. The Product Experience    | The dashboard and screens customers actually use: compliance scores, the assistant, approvals, and the audit explorer.             | A polished product a salesperson can put in front of a prospect.                    | 640 hrs |
| 4. Trust & Hardening         | Independent security testing, speed tuning, verifiable audit exports, and always-on reliability.                                   | An audit-ready product safe to sell to regulated enterprises.                       | 640 hrs |

### Why this sequencing protects the investment

If priorities change, the business already holds demonstrable assets at the end of each stage.

The riskiest, trust-critical work (security testing, reliability) is deliberately scheduled last, once the product is proven.

## 10. Investment & Resourcing

### 10.1 What we are asking to fund

| Item                        | Detail                                                                                   |
|-----------------------------|------------------------------------------------------------------------------------------|
| Engineering effort          | 2,560 hours across four stages (640 hours each)                                          |
| Indicative engineering cost | ~US\$123,000 at the standard blended delivery rate                                       |
| Team                        | An 8-person delivery pod (lead, engineers, design, operations, security/quality)         |
| Indicative build duration   | ~10 weeks of build, plus elapsed time for independent testing and the SOC 2 audit window |

*The engineering figure is an indicative build cost based on the standard internal blended rate and the SRS effort model. It is a planning estimate, not a fixed quote.*

### 10.2 Costs that are additional (flagged honestly)

- Cloud infrastructure and third-party software (ongoing operating cost).
- Independent penetration testing and the SOC 2 audit (specialist vendors).
- Sales, marketing, and customer onboarding to take the product to market.

### What the engineering investment buys

A complete, sellable first version of AuthClaw covering all three capabilities — not a prototype.

An audit-ready posture that lets us sell into regulated enterprises, the highest-value segment.

## 11. Indicative Timeline

Durations below are indicative and assume the 8-person team described above. Independent security testing and the SOC 2 observation window add calendar time that is outside our direct control.

| Milestone               | What it means for the business                 | Cumulative effort |
|-------------------------|------------------------------------------------|-------------------|
| M1 — Foundation ready   | Safe AI traffic + audit recorder demonstrable  | 640 hrs           |
| M2 — Intelligence live  | Leak prevention + compliance assistant working | 1,280 hrs         |
| M3 — Product ready      | Full customer-facing product to demo and sell  | 1,920 hrs         |
| M4 — Audit-ready launch | Hardened, tested, sellable to regulated buyers | 2,560 hrs         |

## 12. Key Risks & How We Manage Them

| Risk (business view)                                   | Likelihood | How we manage it                                                                                    |
|--------------------------------------------------------|------------|-----------------------------------------------------------------------------------------------------|
| A larger competitor moves into the live-gateway space  | Medium     | Move quickly to a sellable v1; focus on the combined wedge none of them currently own.              |
| The product is too slow and frustrates users           | Medium     | Speed is an explicit, tested target; performance tuning is a dedicated stage-4 workstream.          |
| A data leak or breach during development damages trust | Low        | Customer separation and security are built in from stage 1, then independently penetration-tested.  |
| Regulatory expectations shift                          | Medium     | The product is designed to add new frameworks (e.g. EU AI Act, ISO) without rebuilding.             |
| Timeline slips due to external audits                  | High       | Audits are scheduled early and treated as elapsed time, not build time, so they do not surprise us. |

## 13. Success Metrics the Board Can Track

- **Delivery:** each stage completed against its milestone (M1–M4).
- **Product proof:** an independent security test passed and a verifiable audit export demonstrated.
- **Commercial:** first design-partner or paying customer engaged once a sellable version exists (Stage 3+).
- **Trust:** SOC 2 readiness achieved, unlocking the regulated-enterprise segment.

## 14. Strategic Fit

**AuthClaw is a natural extension of the AgentsArchitects portfolio.** Our existing platforms already help enterprises adopt AI; AuthClaw lets them adopt it safely and prove it. It deepens our position in the enterprise AI value chain and adds a recurring-revenue software product to a portfolio aimed at reaching billion-dollar scale before 2030.

It also reinforces our credibility: a company that builds AI compliance software is, by definition, demonstrating compliance leadership to every enterprise buyer it meets.

## Appendix A — Plain-English Glossary

| Term              | What it means                                                                                                                   |
|-------------------|---------------------------------------------------------------------------------------------------------------------------------|
| AI model          | The underlying artificial-intelligence system (e.g. a large language model) that a company sends requests to.                   |
| Gateway           | A checkpoint that all AI traffic passes through, so it can be inspected and made safe.                                          |
| PII / PHI         | Personally Identifiable Information / Protected Health Information — sensitive personal or medical data that must be protected. |
| GDPR              | European data-protection law with significant penalties for mishandling personal data.                                          |
| HIPAA             | US law protecting the privacy and security of health information.                                                               |
| SOC 2             | A widely recognised security and trust standard that customers often require from suppliers.                                    |
| EU AI Act         | New European law regulating AI use, with fines up to €35M or 7% of global turnover for serious breaches.                        |
| Human-in-the-Loop | A safeguard ensuring a person must approve important or risky actions before they happen.                                       |
| Audit trail       | A tamper-proof record of everything that happened, usable as proof for auditors and customers.                                  |

Prepared by Binod Kumar · AgentsArchitects.ai  
<https://agentsarchitects.ai/>